

 Open in New Tab

To print the PDF document, please use the Print button on the toolbar.



Meta Review Status: The meta review and decision have not been released yet.



AI Review from DeepReviewer will be automatically processed



Summary

This paper introduces an adaptive prompt-enhanced score matching method for recovering statistical score functions from partially observed data, primarily under the Missing Completely at Random (MCAR) assumption. The core contribution lies in the integration of a meta-learning

prompt generator that dynamically selects key hyperparameters, such as the Importance-Weighted (IW) sample size, variational inner steps, and learning rates, to optimize the convergence behavior of the model. The authors evaluate their method on synthetic datasets, including multivariate Gaussians, ICA-inspired models, and sparse Gaussian graphical models (GGMs) with star graph structures. The empirical results demonstrate significant improvements in parameter estimation and structure recovery under partial observations. However, the paper's presentation and methodological details are somewhat lacking, and the theoretical underpinnings and comparisons to existing methods are not as robust as they could be. Despite these limitations, the paper addresses a relevant and challenging problem in the field of score matching with missing data, and the adaptive hyperparameter selection mechanism is a novel and interesting approach.

✓ Strengths

The paper makes a valuable contribution to the field of score matching with partially observed data by introducing an adaptive prompt-enhanced method that dynamically selects key hyperparameters. This approach is particularly innovative, as it leverages a meta-learning prompt generator to optimize the convergence behavior of the model across diverse data distributions and missingness patterns. The empirical validation of the method is comprehensive, with experiments conducted on synthetic datasets, including multivariate Gaussians, ICA-inspired models, and sparse Gaussian graphical models (GGMs) with star graph structures. The results show significant improvements in parameter estimation and structure recovery, which is a promising step forward in handling missing data in score matching. The paper also provides a detailed discussion of the bias-variance trade-off inherent in such estimators, which is a crucial aspect of the methodology. The authors' focus on high-dimensional data and the use of zero-imputation for missing values are practical and relevant considerations for real-world applications. Overall, the paper's core idea and empirical findings are strong, and the method has the potential to address a significant gap in the literature on score matching with missing data.

✗ Weaknesses

Despite the paper's promising contributions, several weaknesses and limitations are evident. First, the presentation of the paper is not as clear or structured as it could be. The motivation, methodology, and contributions are scattered and not well-organized, making it difficult to follow the flow of the paper. For instance, the introduction could more explicitly state the gap in existing research that this paper aims to fill, and the methodology section should provide a more detailed, step-by-step explanation of the proposed algorithm, including the meta-learning prompt generator and the mathematical formulation of the score matching objective. The current structure, while containing the necessary sections, does not effectively highlight the novelty and significance of the proposed method (confidence level: high, evidence: poor organization and flow within sections, lack of explicit novelty comparison in the introduction and related work). Second, the novelty of the proposed method is questionable. While the authors claim that their approach is designed for high-dimensional settings, this claim is not adequately substantiated. The method in Givens et al. (2025) appears to be applicable to high-dimensional settings as well, provided that the hyperparameters are tuned appropriately. The authors need to clarify what specific aspects of their method are truly novel and how these aspects provide a significant improvement over existing approaches (confidence level: high, evidence: lack of detailed comparison with Givens et al. (2025) and unsubstantiated high-dimensional claim). Third, the methodological details are inadequate, particularly regarding the adaptive prompt mechanism and the meta-learning process. The paper lacks a comprehensive explanation of how the adaptive prompting works, including the specific algorithms and mathematical formulations involved. Additionally, the paper does not clearly explain how the method addresses the challenges of partially observed data, such as whether it is imputing missing data, working directly with observed data, or using a combination of both (confidence level: high, evidence: conceptual description of meta-learning without algorithmic details, mention of zero-imputation only in experimental setup). Fourth, the paper lacks a detailed theoretical analysis of the proposed method. While the empirical results are promising, a theoretical understanding of the convergence properties, generalization bounds, and the impact of the adaptive hyperparameter selection on the optimization landscape would strengthen the paper. Specifically, the paper does not provide any guarantees on the convergence of the meta-learning prompt generator, nor does it analyze how the dynamic selection of hyperparameters affects the stability and convergence rate of the score matching objective (confidence level: high, evidence: absence of theorems, proofs, or theoretical discussions). Fifth, the paper primarily focuses on the MCAR scenario, which is a significant limitation. The current analysis does not address the potential biases introduced by different missing data mechanisms, such as Missing at Random (MAR) and Missing Not at Random (MNAR). In MNAR settings, the missingness depends on the unobserved values, which could lead to biased parameter estimates and incorrect structure recovery. The paper

should at least discuss the potential challenges and limitations of applying the method to MAR and MNAR data (confidence level: high, evidence: explicit focus on MCAR in the introduction and experimental setup, lack of discussion on MAR/MNAR). Finally, the paper does not provide a detailed comparison with existing methods for handling missing data in score matching. While the authors mention some related work, a more thorough comparison, including a discussion of the advantages and disadvantages of the proposed method compared to alternatives, would be beneficial. Specifically, the paper should compare the proposed method with other imputation-based approaches, as well as methods that directly handle missing data in the score matching objective. A quantitative comparison on benchmark datasets would be necessary to demonstrate the superiority of the proposed method (confidence level: high, evidence: lack of quantitative comparison with other methods in the experimental results).

 Suggestions

To address the identified weaknesses, the paper needs significant restructuring and additional content. First, the authors should reorganize the paper to include clear sections for motivation, problem statement, related work, methodology, experiments, results, and conclusion. The motivation section should clearly articulate the gap in existing research that this paper aims to fill, and the problem statement should precisely define the scope and assumptions of the proposed method, including the type of missingness (e.g., MCAR, MAR, MNAR) and the nature of the data (e.g., continuous, categorical). The related work section should provide a detailed comparison of the proposed method with existing approaches, highlighting the specific limitations of prior work that this paper addresses. The methodology section should include a thorough explanation of the adaptive prompt mechanism and the meta-learning process, with clear mathematical formulations and algorithms. For example, the authors should detail how the prompts are generated and updated, and how the method ensures that the prompts are effective across different data distributions and missingness patterns. The experiments section should clearly describe the experimental setup, including the datasets used, the evaluation metrics, and the baselines. The results section should present the findings in a clear and concise manner, with appropriate statistical analysis and visualizations. The conclusion should summarize the key contributions of the paper and discuss potential future directions. Second, the authors need to provide a more precise definition of what makes their method novel. If the claim of being

 Questions

1. Can the authors provide a detailed algorithmic description of the meta-learning prompt generator, including the specific algorithms and mathematical formulations used to generate and update the prompts? How does the method ensure that the prompts are effective across different data distributions and missingness patterns? 2. What are the theoretical guarantees for the proposed method, particularly regarding the convergence of the meta-learning prompt generator and the impact of adaptive hyperparameter selection on the optimization landscape? How do these guarantees compare to existing methods for score matching with missing data? 3. How does the proposed method handle the challenges of score matching with partially observed data, specifically in terms of imputation, direct observation, or a combination of both? Can the authors provide a more detailed explanation of the zero-imputation strategy and its implications for the method's performance? 4. Can the authors extend their analysis to more complex missing data mechanisms, such as MAR and MNAR? What are the potential biases introduced by these mechanisms, and how can the proposed method be adapted to handle them? What are the limitations of the current method in these scenarios, and what are the potential solutions for future research? 5. How does the proposed method compare to existing methods for handling missing data in score matching, both in terms of performance and computational complexity? Can the authors provide a detailed quantitative comparison on benchmark datasets, including a discussion of the advantages and disadvantages of their method compared to alternatives?

 Scores

Soundness: 1.75

Presentation: 1.25

Contribution: 1.5

Rating: 3.5

 AI Review from ZGCA will be automatically processed Summary

The paper studies score function recovery from partially observed data under Missing Completely at Random (MCAR, ~30% missing). It compares marginal Importance-Weighted (Marg-IW) and marginal Variational (Marg-Var) schemes and introduces a 'meta-learning prompt generator' that adaptively selects key hyperparameters (e.g., $r \in \{5, 10, 50\}$, $L \in \{1, 5, 10\}$, learning rates, truncation τ , and projection count) to stabilize training. The score model is Gaussian $s_\theta(x) = -P_\theta(x - \mu_\theta)$ with $P_\theta = L_\theta L_\theta^T$ (diagonal exp for PD), optimized using a surrogate MSE to the true score $s_{\text{true}}(x) = -P_{\text{true}}(x - \mu_{\text{true}})$, restricted to observed entries via masks M (e.g., Eqns. $L_{\text{obs}}(\theta)$ and $L_{\text{IW}}(\theta)$). Stabilization uses log-sum-exp and gradient clipping; GGM experiments add an L1 penalty on off-diagonal precision entries. Experiments on synthetic Gaussians and a 10D GGM star-graph (MCAR 30%) report decreasing surrogate loss ($9.687 \rightarrow 0.094$) and AUC improvements ($0.219 \rightarrow 0.972$).

 Strengths

- Addresses a relevant problem: score estimation under partial observations with MCAR masking, explicitly formulating masked losses (e.g., $L_{\text{obs}}(\theta)$ with $M \odot [\cdot]$).
- Clear Gaussian score parameterization with PD constraints ($P_\theta = L_\theta L_\theta^T$; exp on diag), plus practical stabilizers (gradient clipping; mentioned log-sum-exp).
- Explores both Marg-IW and Marg-Var variants and articulates the bias-variance considerations in missing data settings.
- Promising structural recovery in GGM star graphs (ROC AUC rising to 0.972), and convergence trends reported for the Gaussian case.
- Recognizes limitations (e.g., zero-imputation bias, MCAR-only focus) and proposes future extensions to MAR/MNAR and diffusion-based score models.

 Weaknesses

- Central contribution (the 'meta-learning prompt generator') lacks algorithmic detail: no explicit selection criterion, update rule, or objective beyond 'convergence behavior and gradient statistics'; makes the key novelty hard to assess or reproduce (Section 3: 'updated on the fly' without specification).
- No quantitative comparisons to state-of-the-art baselines (e.g., 'Score Matching with Missing Data' IW/Var, robust SM). The paper contains comparison tables but no empirical benchmarking against prior methods.
- Evaluation is narrowly scoped to 30% MCAR; does not explore different missingness rates or MAR/MNAR, limiting generalizability.
- The training objective uses a surrogate MSE to s_{true} (e.g., Eqns. $L(\theta)$, $L_{\text{obs}}(\theta)$), which requires access to ground-truth parameters; this is non-standard for score matching and limits applicability beyond synthetic settings (and is not viable for CIFAR-10).
- Inconsistency between the experimental setup (Section 4 lists a CIFAR-10 subset) and presented results (Section 5 reports only Gaussian/GGM; no CIFAR results).
- Methodological gaps: (i) IW estimator not fully specified (weights, proposal, variance control), (ii) variational inner-loop objective and reparameterization details omitted, (iii) 'sliced score matching' mentioned but not defined, (iv) L1 penalty strength for GGM not reported.
- Reported metrics are point estimates without variance/confidence intervals or multiple runs; claimed ablations are not shown. Parameter error remains relatively high ($3.033 \rightarrow \sim 2.030$) despite large loss reductions, and is not analyzed.
- Clarity/reproducibility concerns: heavy repetition, limited derivations, and no pseudo-code for the adaptive component; zero-imputation is used to form tensors (Section 4), which can bias results without sensitivity analysis.

? Questions

- Please specify the 'meta-learning prompt generator' precisely: What signal(s) are monitored (loss slope, gradient norms, validation metric)? What is the selection rule (e.g., bandit, RL, Bayesian optimization, threshold heuristics)? How often are hyperparameters updated, and is there any backpropagation through the selection?
- For Marg-IW: What is the imputation proposal distribution? How are importance weights computed and normalized? Is there any variance reduction (e.g., control variates)? How is r chosen adaptively beyond sampling from $\{5, 10, 50\}$?
- For Marg-Var: What is the explicit variational objective and family? What are the inner-loop update equations for L steps, and how are $\eta\phi$ and L adapted? Do you backprop through inner steps (i.e., bi-level optimization)?
- Sliced score matching: Please define the slice projections and the exact loss you used; how does proj_count enter the estimator, and where is it combined with the masked loss?
- GGM sparsity: What is the L1 penalty coefficient (λ) and how was it tuned (fixed vs adaptive)? How sensitive are AUC results to λ ?
- Baselines: Can you provide quantitative comparisons against (i) fixed-hyperparameter Marg-IW/Var, (ii) robust score matching variants, and (iii) standard graphical lasso under MCAR with imputation? This will contextualize the reported AUC gains.
- Robustness: How does performance vary across missingness rates (e.g., 10%, 30%, 50%) and different missingness mechanisms (MAR, synthetic MNAR)?
- CIFAR-10: You include a CIFAR-10 subset in the setup; please provide results or clarify why it was omitted. How is s_{true} defined in that setting given the Gaussian score model?
- Statistical reporting: Can you report mean $\pm$ std over multiple runs and include ablations isolating the contribution of the adaptive component vs fixed grids or random search?
- Compute and complexity: What is the runtime and memory cost of adaptive selection compared to fixed hyperparameters? How does cost scale with d , r , and L ?

! Limitations

- Relies on a surrogate loss to s_{true} (e.g., $L(\theta)$, $L_{\text{obs}}(\theta)$), which requires ground-truth parameters and is not available in real data; limits external validity.
- Evaluation restricted to 30% MCAR; no experiments varying missingness rate or addressing MAR/MNAR, which are common in practice.
- Zero-imputation for tensor formation (Section 4) can bias estimation; no sensitivity analysis or comparison to alternatives (mean/mask embedding/learned imputation).
- Lack of baselines, variance estimates, and ablations reduces confidence in the claimed gains of the adaptive component.
- Model class is linear Gaussian scores ($s_{\theta}(x) = -P_{\theta}(x - \mu_{\theta})$); mismatch with non-Gaussian settings (e.g., CIFAR-10) is not addressed.
- Potential negative impact: If deployed naively on sensitive domains (e.g., healthcare) with non-MCAR missingness, the combination of zero-imputation and uncalibrated adaptivity could introduce biased inferences, affecting downstream decisions.

Image Evaluation

Cross-Modal Consistency: 28/50

Textual Logical Soundness: 18/30

Visual Aesthetics & Clarity: 16/20

Overall Score: 62/100

Detailed Evaluation (≤ 500 words):

Image-First Understanding (visual ground truth)

- Figure 1/(a): Line plot with legend “Score Matching Loss” (blue) and “Parameter Error” (orange), y-axis “Loss/Error”, x-axis “Iteration” (0–300). Both curves monotonically decrease; orange appears near zero by the end.
- Figure 1/(b): Dual-axis plot. Blue left-axis “GGM Score Matching Loss” rapidly decays; red right-axis “ROC AUC” shows markers at ~50, 100, 150, 200 iterations increasing to ≈ 0.95 .
- Figure-level synopsis: (a) Gaussian training convergence; (b) GGM training and structure-recovery metric; complementary quantitative views across two settings.

1. Cross-Modal Consistency

- Major 1: Table mismatch – text says it summarizes missingness mechanisms but shows hyperparameters. Evidence: Sec 1 “Table below summarizes various missingness mechanisms” vs hyperparameter table.
- Major 2: Parameter-error numbers conflict with Fig. 1(a). Evidence: Sec 5 “parameter error... 3.033... to about 2.030” vs Fig. 1(a) orange approaches ~ 0 by 300.
- Major 3: Claimed ROC AUC at 250/300 not visible in figure. Evidence: Sec 5 “0.972 at both iterations 250 and 300” vs Fig. 1(b) shows red markers only at 50,100,150,200.
- Major 4: CIFAR-10 dataset listed but no corresponding results. Evidence: Sec 4 describes CIFAR-10 (3072-d) yet no figure/table reports CIFAR outcomes.
- Minor 1: Duplicate comparison tables (Abstract, Sec 2) with slightly different attributions (Schwank vs Givens). Evidence: Sec 2 table header vs References.
- Minor 2: Sub-figure labels (a)/(b) are in text, not on panels. Evidence: Figure 1 panels lack internal (a)/(b) tags.

2. Text Logic

- Major 1: Use of s_{true} with real CIFAR data is unclear (true P_{true} , μ_{true} unknown). Evidence: Sec 4 includes CIFAR but s_{true} defined via ground-truth parameters (Sec 1/3).
- Minor 1: Claim that formulation “inherently accounts for missingness” conflicts with later zero-imputation note. Evidence: Sec 2 “inherently accounts...” vs Sec 4 “use of zero-imputation”.
- Minor 2: Author misattribution across sections (Schwank vs Givens). Evidence: Sec 2 vs References.

3. Figure Quality

- Minor 1: Fig. 1(b) dual-axis is clear but lacks tick labels for ROC AUC points beyond 200; add markers for 250/300. Evidence: Fig. 1(b).
- Minor 2: Embed (a)/(b) tags on panels and add numeric callouts at iteration 50 and 300 for easier verification. Evidence: Figure 1.

Key strengths:

- Clear, modular method description; positive-definite parameterization and stabilization choices are reasonable.
- GGM plot conveys simultaneous loss and ROC AUC trends effectively.

Key weaknesses:

- Multiple figure–text inconsistencies on key quantitative claims.
- Missing CIFAR results despite being a core dataset.
- Ambiguity about ground-truth scores on real data; table mislabeling reduces trust.



Scores

Originality: 2

Quality: 2

Clarity: 2

Significance: 2

Soundness: 2

Presentation: 2

Contribution: 2

Rating: 4 AI Review from SafeReviewer will be automatically processed **Summary**

The paper proposes a method to learn score functions from data with missing entries, using a meta-learning prompt generator to dynamically select key hyperparameters. The method is evaluated on synthetic datasets, including sparse Gaussian Graphical Models (GGMs), and shows improved performance in parameter estimation and structure recovery compared to traditional methods.

****Soundness:**** 2.0

****Presentation:**** 1.0

****Contribution:**** 1.67

 Strengths

The paper addresses a significant and relevant problem in machine learning: the challenge of learning score functions from partially observed data. This is a crucial issue in many real-world applications where data is often incomplete. The core idea of using a meta-learning prompt generator to dynamically select hyperparameters is innovative and has the potential to improve the performance and robustness of score matching methods. The authors' attempt to adapt score matching to handle missing data, particularly through the use of marginal importance-weighted and marginal variational approaches, is a valuable contribution. The experimental results, while limited to synthetic data, do show promising improvements in surrogate loss reduction, parameter estimation accuracy, and structural recovery, particularly in the Gaussian Graphical Model setting. The inclusion of an L1 penalty for sparsity in the GGM experiments is a sensible choice and contributes to the improved performance. The paper also attempts to address the numerical stability of the method through the use of log-sum-exp stabilization and gradient clipping, which are important considerations when dealing with score matching. The authors' recognition of the limitations of existing methods for handling missing data and their attempt to address these limitations through an adaptive approach is a positive aspect of the work. The paper also highlights the potential for future work, such as extending the method to handle more complex missing data mechanisms and integrating diffusion-based denoising models. These potential extensions suggest that the proposed method could be a valuable contribution to the field if further developed and refined.

 Weaknesses

After a thorough review of the paper and the reviewer comments, I've identified several significant weaknesses that undermine the paper's current presentation and conclusions. First and foremost, the paper suffers from a lack of clarity and precision in its writing. The abstract

introduces notation, such as μ and P , without prior definition, making it difficult to understand for readers unfamiliar with the specific context. This issue extends throughout the paper, with undefined terms and inconsistent notation, creating a barrier to comprehension. For example, the term 'prompt' is used extensively without a clear definition, leaving the reader to guess at its meaning and implementation. This lack of clarity is particularly problematic given the central role of the meta-learning prompt generator in the proposed method. The paper also lacks a clear explanation of the meta-learning process. The architecture and training procedure of the prompt generator are not described in sufficient detail, making it difficult to understand how it functions and how it contributes to the overall performance of the method. The paper mentions that the prompt generator selects hyperparameters from a predefined set, but the criteria for selection and the mechanism of the generator are not explained. This lack of transparency makes it difficult to assess the novelty and effectiveness of the proposed approach. Furthermore, the experimental section is not well-structured and lacks sufficient detail. The experimental setup is described in a way that combines the setup for all experiments, making it difficult to distinguish between the different experimental scenarios. The results are presented in a dense and somewhat confusing manner, with tables and text interwoven. The paper also lacks a clear explanation of the baselines used for comparison. While the authors state that they compare against classical score matching methods and naive imputation techniques, these baselines are not explicitly defined or described, making it difficult to evaluate the relative performance of the proposed method. The paper also lacks a thorough analysis of the experimental results. The discussion of the results is primarily descriptive, focusing on the observed improvements without delving into the underlying reasons for these improvements. The paper also lacks a discussion of the limitations of the proposed method and the potential for future work. While the discussion section touches on future directions, it does not explicitly address the limitations of the current approach. Finally, the paper lacks a rigorous theoretical analysis of the proposed method. The paper does not provide any theoretical guarantees or analysis of the method's properties, such as convergence rates or generalization bounds. This lack of theoretical justification makes it difficult to assess the reliability and robustness of the method. The paper also does not adequately address the potential bias introduced by the zero-imputation strategy used for handling missing data. While the authors acknowledge the use of zero-imputation, they do not discuss its potential impact on the accuracy of the estimated score function. The paper also lacks a clear explanation of the role of the L1 penalty in the Gaussian Graphical Model experiments. While the authors state that the L1 penalty is used to enforce sparsity, they do not provide a detailed explanation of how it contributes to the improved performance. In summary, the paper suffers from a lack of clarity, insufficient methodological detail, inadequate experimental validation, and a lack of rigorous theoretical justification. These weaknesses significantly undermine the paper's current presentation and conclusions, making it difficult to assess the practical utility and generalizability of the proposed method. The confidence level for these weaknesses is high, as they are consistently supported by the paper's content and the reviewers' comments.

Suggestions

To address the identified weaknesses, I recommend several concrete and actionable improvements. First, the authors must significantly improve the clarity and precision of their writing. This includes providing clear definitions for all terms and notations, especially in the abstract and when introducing new concepts. The authors should also ensure consistent notation throughout the paper and avoid using undefined terms. The explanation of the meta-learning prompt generator needs to be significantly enhanced. The authors should provide a detailed description of its architecture, including the type of model used (e.g., neural network, gradient-based method), the input and output spaces, and the training procedure. The paper should also clearly explain the mechanism by which the prompt generator selects hyperparameters and the criteria used for selection. The experimental section needs to be restructured and clarified. The authors should separate the experimental setup for each dataset and clearly state the experimental objectives. The results should be presented in a clear and concise manner, with tables and figures that are easy to understand. The baselines used for comparison should be explicitly defined and described, and the authors should provide a detailed analysis of the experimental results, explaining the reasons for the observed improvements and discussing the limitations of the proposed method. The authors should also include a more thorough theoretical analysis of the proposed method. This should include an analysis of the method's properties, such as convergence rates and generalization bounds. The authors should also address the potential bias introduced by the zero-imputation strategy and discuss alternative imputation techniques. The role of the L1 penalty in the Gaussian Graphical Model experiments should also be clarified, with a detailed explanation of how it contributes to the improved performance. The authors should also consider including more complex datasets in their experiments, such as CIFAR-10 or ImageNet, to demonstrate the generalizability of their method. The paper should also include a more thorough discussion of the limitations of the proposed method and the potential for future work. This should include a discussion of the assumptions made by the method and the potential for extending it to handle more complex missing data mechanisms and more complex data distributions. Finally, the authors should carefully proofread the paper to correct any grammatical errors and ensure that the paper is well-organized and easy to follow. By addressing these issues, the authors can significantly improve the

quality and impact of their work. These suggestions are directly tied to the identified weaknesses and aim to provide concrete steps for improvement. The scope of these changes is realistic and achievable within the context of a revised paper.

? Questions

Several key uncertainties and methodological choices remain unclear after my review of the paper. First, I am curious about the specific architecture and training procedure of the meta-learning prompt generator. What type of model is used for the generator? How is it trained, and what is the loss function used for training? How does the generator explore the hyperparameter space, and what criteria are used to select the optimal hyperparameters? Second, I am interested in the theoretical properties of the proposed method. What are the convergence rates of the method, and what are the conditions under which it is guaranteed to converge to the true score function? What is the impact of the missing data mechanism on the convergence of the method? Third, I would like to understand the limitations of the zero-imputation strategy used for handling missing data. How does this strategy affect the accuracy of the estimated score function, and what are the potential biases introduced by this approach? Are there alternative imputation techniques that could be used to mitigate these biases? Fourth, I am curious about the role of the L1 penalty in the Gaussian Graphical Model experiments. How does the L1 penalty contribute to the improved performance, and what is the impact of the regularization parameter on the results? Fifth, I would like to know more about the generalizability of the proposed method. How does the method perform on more complex datasets, such as CIFAR-10 or ImageNet? What are the challenges of applying the method to these datasets, and how can these challenges be addressed? Finally, I am interested in the potential for extending the proposed method to handle more complex missing data mechanisms, such as MAR or MNAR. What are the challenges of extending the method to these mechanisms, and what are the potential solutions? These questions target core methodological choices and seek clarification of critical assumptions, aiming to provide a deeper understanding of the proposed method and its limitations.

 Scores

Soundness: 2.0

Presentation: 1.0

Contribution: 1.67

Rating: 3.0

Keywords

Extract Keywords

Click the button to extract keywords

Insights

Extract Insights

Click the button to extract insights